



● Llama-8B    ■ OLMoE-1B-7B    ▲ Qwen-7B    ◆ Qwen-4B


BoolQ




LogiQA



MMLU



OpenBookQA



Request rate (requests/s)

p95 inference latency (ms)